

TDT4225 Assignment 2 – Kristoffer Dalby

Exercise 1 – I/O Programming

Code:

```
import os
import array
import time

BLOCKSIZE = 8192
NBLOCKS = 131072

# Define a array of 8192 bits
BLOCK = array.array('b', [1 for x in range(BLOCKSIZE - 64)])

def pretty(ms, size):
    '''
    This method pretty prints the output
    '''
    megabyte = size * 1024
    throughput = megabyte / (ms / 1000)
    throughput = '{0:.2f}'.format(throughput)
    ms = '{0:.2f}'.format(ms)
    print('{} GB\t{} MB/s\t{} ms'.format(size, throughput, ms))

print('{}\t{}\t{}'.format('EXT4', 'Throughput', 'Time'))

for size in [1, 2, 4, 8, 16, 32]:
    # Open the file in binary form
    with open('test{}'.format(size), 'wb') as f:

        start = time.clock()

        for i in range(NBLOCKS * size):
            f.write(BLOCK)

        # Force write to file
        f.flush()
        # Force write to disk
        os.fsync(f.fileno())
        os.system('sync')
        stop = time.clock()

        interval = stop - start

        pretty(interval*1000, size)
```

The file is also uploaded to its learning.

Result:

Device: iMac
OS: Mac OS X 10.10
Filesystem: HFS+
Disk: SSD

HFS+	Throughput	Time
1 GB	679.58 MB/s	1506.82 ms
2 GB	681.96 MB/s	3003.11 ms
4 GB	683.43 MB/s	5993.29 ms
8 GB	655.67 MB/s	12494.10 ms
16 GB	663.56 MB/s	24691.05 ms
32 GB	658.44 MB/s	49765.79 ms

Device: Lenovo X220i
OS: Debian 8 Jessie
Filesystem: EXT4
Disk: HDD

EXT4	Throughput	Time
1 GB	553.62 MB/s	1849.63 ms
2 GB	466.81 MB/s	4387.19 ms
4 GB	422.20 MB/s	9701.66 ms
8 GB	411.10 MB/s	19926.89 ms
16 GB	377.90 MB/s	43355.18 ms
32 GB	368.84 MB/s	88840.23 m

Device: VM
OS: Debian 8 Jessie
Filesystem: EXT4
Disk: HDD 10k RPM in Raid 50

EXT4	Throughput	Time
1 GB	533.68 MB/s	1918.76 ms
2 GB	530.99 MB/s	3856.94 ms
4 GB	521.46 MB/s	7854.93 ms

This server only has 10GB of space, with 2GB occupied.

I think the implementation is correct, and I have verified the size of each file to be the correct size. I personally think that it may seem a little to fast, but it may be that the program is only writing 1s.

Exercise 2 – Various question on file systems

a)

S5FS with 1KB datablock can have files up to 2 GB as a file needs to fit inside a unsigned integer. In theory the maximum size is 16GB, but you cannot achieve it.

b)

When S5FS frees and allocates inodes it happens less often than datablocks. This means that having them on the disk is not that problematic. The supernode caches all the nodes so finding them can go relatively fast to. The technique of having the inodes on the disk also makes the system more crash tolerant as the inodes will not be damaged.

c)

The S5FS file system generally uses small block sizes which means that data that belongs together can be spread out all over the disk when the data is defragmented. FFS uses larger blocks which means that data will not be spread that much around on the disk. FFS also has a technique called block interleaving, this makes the blocks not appear sequentially after each other. This improves performance on disks where the disk cannot read or write fast enough so that the disk needs to take a whole turn around before the next read or write.

d)

Journaling in filesystems works like in databases, it keeps a journal of every read and write operation. This means that the system can recreate, and/or rollback operations that did not go correctly or got aborted by for example a crash.

Soft updates writes the order of what has happened to volatile storage(e. g. RAM). This means that if the computer crashes or loses power the writes in the cache does not have any effect and will not or have not happened.

e)

Cylinder groups is a concept where a cylinder group consists of a collection of inodes and data blocks. The technique is based on the idea that data that belong together is put in the same cylinder. This includes both data and metadata. This can reduce fragmentation as data can be put in sequence on the disk. Seek time and latency can be reduced for the same reason.

Extent-based block allocation is based mainly on the idea that if you read a file, you will likely read files that are close to it afterwards. An extent can be looked at as a contiguous area of storage reserved for a file in the file system. A file can consist of zero or more extents. Extent will allow metadata for big files to be saved smarter, but it will also cause a bad sector error to be far more damaging as it can corrupt metadata.

f)

RAID is a technique used to split data over multiple storage devices. It can both increase performance or redundancy and sometimes both.

RAID 0

Data is spliced evenly across the drives in the RAID set. There is no redundancy and the performance is usually increased. If one drive fails, all data is lost.

RAID 1

Mirrors all data to every drive in the RAID set. There is full redundancy but the performance may suffer. RAID1 is usually fast at reads as it can be done by either drive. Write speed is usually equal to the write rate of the slowest disk.

RAID 2

Stripes the data at the bit level. The disks are generally synchronized in a way that makes the associated data reside on the same place on each disk. RAID 2 cannot generally service multiple requests simultaneously, but can hit extremely high transfer rates as data from disks that belong together can be read at the same time.

RAID 3

Uses a dedicated parity disk for error correction. The data is striping on the byte-level. As in RAID 2, this implementation collects data and synchronize the hard drive so that long sequential reads and writes can go very fast.

RAID 4

Uses a dedicated parity disk. Uses block-level striping.

RAID 5

Mostly replaces RAID 3 and 4.

It uses block-level striping, but the parity is distributed over all the disks. Upon failure, the reads can be calculated based on the parity so no data is lost.

RAID 6

Extends RAID 5 with another parity block

RAID 10 (1 + 0)

Combination of RAID 1 and RAID 0

RAID 50 (5 + 0)

Combination of RAID 5 and RAID 0

RAID 60 (6 + 0)

Combination of RAID 6 and RAID 0